

Academic Application Tracker

UX Field Study Report

Three-month simulated usage across three personas
Sept – Nov 2026
Fact-verified against the v0.14.0-dev codebase

Report date	2026-05-22
Codebase reference	feat/ui-redesign-v0.14.0 @ commit 2ded3ba
Audience	Marketing · Product · Engineering
Length	20 pages, 6 figures, 7 tables
Classification	Internal — pre-decision research

Table of Contents

1.	Executive Summary	3
2.	Study Methodology	4
3.	Personas at a Glance	5
4.	Persona 1 — Aisha K., life-sci postdoc	6
5.	Persona 2 — James M., humanities TT applicant	7
6.	Persona 3 — Wei L., CS PhD mixed search	8
7.	Cross-cutting Friction Themes	9
8.	Friction Severity Heatmap	10
9.	Verified Bug-class Findings	11
10.	Churn-risk Timeline + Retained Scope	12
11.	What's Working — Strengths to Preserve	13
12.	Competitive Landscape Analysis	14
13.	Competitive Feature Matrix	15
14.	Prioritized Recommendations (P0)	16
15.	Prioritized Recommendations (P1, P2) + Effort-Impact	17
16.	Strategic Roadmap & Audience-specific Takeaways	18
17.	Appendix A — Verified Code References	19
18.	Appendix B — Methodology Limitations + References	20

1. Executive Summary

Three independent persona simulations (a life-sciences postdoc, a humanities tenure-track applicant, and a computer-science PhD running a mixed academic + industry search) used the Academic Application Tracker (AAT) for a simulated three-month cycle, Sept–Nov 2026. Each persona's journey was traced against the actual v0.14.0-dev codebase; the friction points, near-quit moments, and workarounds reported below are tied to specific file and line numbers and have been spot-verified against the source.

Headline finding.

The dominant churn risk is structural, not cosmetic. The **recommender entity is modelled per-position** rather than as a global first-class entity, which forces $N \times M$ data entry on Day 1. Both the postdoc and the humanities applicant nearly abandoned the tool within the first hour; both permanently demoted the Recommenders page to a manual filing-cabinet workflow.

Secondary themes (all three personas).

- Hardcoded vocabulary in *config.py* requires Python edits to extend (blocks non-coders entirely; converts a power user into a contributor).
- No bulk operations — single-row select everywhere; one-recommender-per-row submission updates; one-cell-per-edit requirement toggles.
- Cascade asymmetries: **R2 promotes on add_interview but does not retract on delete_interview** (real defect, verified at database.py:707–719).
- Mobile/tablet experience is broken; lost $\approx 30\%$ of one persona's tracking sessions to layout collapse under *layout="wide"*.

What's working (preserve at all costs).

- The editorial-brutalist visual identity (*ui.py*) drew unprompted positive reactions from all three personas — rare for a research-internal Streamlit app.
- The Upcoming-deadlines panel + urgency banding is the single feature that *genuinely outperforms a spreadsheet* for two of three personas.
- Markdown exports on every write (DESIGN D4) build durable trust that the data survives the app itself.
- The *config.py* import-time invariants converted a near-fork into an upstream pull request from the power-user persona — the highest praise a codebase can earn from a developer-evaluator.

Verdict.

All three personas retained the tool through Month 3 — none churned outright — but two operate at roughly 50% of intended scope, maintaining parallel Google Docs / Sheets for the workflows AAT could not absorb. The recommender-entity refactor (P0) is the single biggest churn-reduction lever; shipping it before any new feature work will lift retention from *de facto* half-use to genuine primary-workflow status.

2. Study Methodology

Design.

Simulated longitudinal field study with three composite personas, each constructed from common applicant archetypes in their field. Each persona's three-month timeline was synthesized by an independent reviewer reading the v0.14.0-dev source in full (DESIGN.md, GUIDELINES.md, app.py, pages/*.py, ui.py, config.py, database.py, docs/dev-notes/self-host-setup.md) and walking through the page-by-page workflows for the persona's most likely operational sequence.

Reviewer brief.

Each reviewer was instructed to be a tough, honest user — not a friendly tester — and to surface at least eight concrete frictions and at least two near-quit moments per persona, citing file and line references so the engineering team could locate root causes.

Verification.

Every load-bearing file:line citation in this report was spot-checked against the current source before publication. Two claims required reclassification:

- Persona 1's report that *applied_date* does not promote status to *[APPLIED]* is incorrect at the data-layer: the R1 cascade is implemented at **database.py:506–510**. Reclassified as a *UI feedback gap* (the promotion happens silently from the user's perspective — no toast on the Applications page) rather than a logic bug.
- Persona 3's report of an R2 asymmetry on *delete_interview* is confirmed: **database.py:707–719** deletes the interview row but has no symmetric *[INTERVIEW]* → *[APPLIED]* retraction. Filed as a real defect.

Scope limits.

This is a simulated study, not a live human study. Findings about visual appeal, friction intensity, and trust shake are reasoned from realistic usage patterns but have not been validated against real user telemetry or interview data. The competitive analysis is desk-research only.

Reproducibility.

Charts and the full PDF are regenerated by `.venv/bin/python scripts/build_ux_report.py`. All friction theme scores, competitive matrix cells, and effort-impact placements live in that script's data arrays; revising any data point only requires editing the script and re-running.

3. Personas at a Glance

Axis	P1 — Aisha K.	P2 — James M.	P3 — Wei L.
Field	Computational biology	Comparative literature	Machine learning / NLP
Cycle target	Fall 2026 postdoc	Fall 2027 TT faculty	Postdoc + industry research
Tracked positions (3 mo)	~40	~22	~51 (15 ac. + 30 ind. + 6 other)
Recommenders	4	5	3
Tech comfort	High (Python dev)	Moderate (no Python)	Very high (would fork)
Primary device	MacBook + iPad on bus	Saturday desktop sessions	Linux desktop
Near-quit moments	2	2	2 near-fork (resolved)
Outcome (Month 3)	Stays w/ caveats (B-)	Stays at ~50% scope	Stays + upstream PR




Figure 1. Persona profile across five normalized axes.

4. Persona 1 — Aisha K., life-sciences postdoc applicant

Profile. Fourth-year PhD in computational biology at a US R1; targeting ~40 fall 2026 postdoc positions over a three-month cycle. Four recommenders; high Python fluency; tracks on a MacBook at her desk and on an iPad on the bus.

Month-1 onboarding.

Setup is smooth — clone, venv, *streamlit run app.py*. The dashboard's editorial design lands well. Quick-add covers six fields. Within 40 minutes she has logged 12 positions. Two hours in she discovers the **schema-UI gap**: *location*, *source*, *mentor*, *point_of_contact*, *portal_url*, *stipend* and four other columns exist in the *positions* table (database.py:77–94) but are written nowhere in the UI. Quote: “*They built half the schema and then forgot to wire it up?*”

Day 14 — near-quit #1.

Adding her first recommender, she discovers the per-position model (database.py:135; pages/3_Recommenders.py:311): 4 recommenders × 40 positions = 160 manual add operations. She stops at row 15.

“This is a database UI, not a job tracker. I’m doing data entry, not application strategy. I came here to stop using a spreadsheet, not to build a worse one.”

Resolution: she enters recommenders only for her top-15 positions; the dashboard alert panel becomes silently incomplete for the remaining 25.

Mid-cycle frictions.

- iPad layout collapses under *layout="wide"*; ~30% of her intended bus-commute tracking sessions are abandoned.
- *Compose Reminder Email* opens a *mailto:* with an empty **To** field — the *recommenders* schema has no email column (verified: database.py:135–148, pages/3_Recommenders.py:142).
- Setting *applied_date* on the Applications page does promote status via R1 (database.py:506–510) but the page emits no visible feedback, so the promotion feels invisible.

Nov 18 — near-quit #2.

36 hours before the Stanford deadline, a fast tab-switch makes the Status selectbox flicker stale state, briefly showing *[SAVED]* for a row she'd already promoted to *[APPLIED]*. She refreshes; the state is correct. But the trust hit is real. “*If this is wrong, my whole pipeline view is wrong.*” She stays out of pure inertia, not loyalty.

Month-3 verdict.

Still using AAT on her laptop, but has **stopped opening it on the bus** and maintains a parallel Google Sheet for recommender coverage. The tool has become “where I write the canonical state once, then never edit from a touch device.” Grade: B–.

5. Persona 2 — James M., humanities tenure-track applicant

Profile. Third-year postdoc in comparative literature; targeting ~22 TT positions for a fall 2027 start. Five recommenders (almost every faculty application asks for all five); 8–11 separate documents per application including a heavily customised cover letter and a 30-page writing sample whose chapter varies per posting. Moderate tech comfort — will not edit Python.

Day-1 setup.

Successful clone + install + *streamlit run*. The editorial design lands as charming rather than fussy. He enters nine positions through Quick-Add over a Saturday morning. The Requirements tab forces a vocabulary gap: there is no *Teaching Portfolio*, no *Job Talk Sample*, no *Syllabi* in *REQUIREMENT_DOCS* (config.py:286–295). He pushes them into the Notes freetext, knowing the dashboard's Materials Readiness will be lying.

Day 1 — near-quit #1.

Adding his five recommenders × 22 positions = 110 manual rows. He stops at row 15. Demotes the Recommenders page on Day 1 from *workflow* to *filing cabinet*. Tracks his real letter-writer state in his existing Google Doc.

Month 2 — the materials-versioning gap becomes the story.

Hopkins wants 25 pages of the Borges chapter; UPenn wants 35 pages of a different (Cortázar) chapter; Cornell wants the Borges chapter with a job-talk-framed rewrite. The Materials tab tracks only a boolean: *done_writing_sample* = 1 = “I sent something.” It records *nothing* about which chapter, which version, or which file. Quote: “*The thing the tool helps me with — am I done? — isn't what I most need to track.*”

Oct 31 — near-quit #2.

Misses the Yale 5pm submission window by minutes. The urgency pill said *T-7D* all week; his Google Calendar's 7-day-out reminder fired the same signal. AAT added zero independent value to deadline awareness for that application. Stays because switching tools 36 hours from a deadline is its own risk.

Month 3 — settled but narrowed.

Uses Dashboard + Opportunities + Applications. Recommenders page effectively abandoned. Keeps Google Docs for cover-letter versioning, writing-sample-per-department tracking, and per-interview prep notes. Saves a Dropbox copy of *postdoc.db* manually on Nov 22 after realizing the markdown export does not protect his state.

Verdict.

Stays at ~50% scope. Would recommend with reservations. Would not pay for it.

6. Persona 3 — Wei L., CS PhD running a mixed search

Profile. Fifth-year PhD, ML / NLP. Mixed search: ~15 academic positions (postdoc + TT) and ~30 industry research roles (DeepMind, Anthropic, FAIR, smaller labs, two YC AI startups). Three recommenders. Very high tech comfort — would read the source and consider forking. Currently on a Notion database.

Day-1 onboarding.

Reads *DESIGN.md* and *config.py* end-to-end before adding data. Quote: “OK, fine, this is a real codebase.” The four-layer architecture and import-time invariants earn explicit respect.

Day 7 — near-fork #1.

Adds his first recruiter screen as an interview row. R2 cascade (*database.py*:629) immediately promotes the position to *[INTERVIEW]*. A 25-minute recruiter chat is now visually equivalent to a Stanford committee onsite on the dashboard funnel. Considers forking; instead opens a GitHub issue draft. **Resolved by codebase extensibility** — the *config.py* invariants make extension feasible.

Day 28 — near-fork #2.

Maintains a personal branch with new statuses (*[GHOSTED]*), new interview formats (Recruiter Screen, Tech Screen, Onsite Loop), new response types, and a CSV importer. Reads the v2 vision in *roadmap.md* and files three GitHub issues instead of forking. Eventually lands one upstream PR (recruiter-screen R2 exemption).

Nov 4 — real defect filed.

Deletes the only interview row on a position that had been auto-promoted to *[INTERVIEW]*. Status does not retract. Files an issue: “R2 promotion is not symmetric with *delete_interview* — orphan *[INTERVIEW]* status after deleting the only interview row.” Verified against *database.py*:707–719: the delete path has no reverse cascade.

Verdict.

Daily use through Month 3. **Contributes upstream** rather than forking — the highest-praise outcome a codebase can earn from a developer-evaluator. Net cost: he maintains a personal branch for industry-specific extensions the upstream hasn't yet absorbed.

7. Cross-cutting Friction Themes

Fifteen friction themes were identified across the three persona studies. The table below ranks them by composite churn-risk weight (frequency × severity × persona reach). Cells use **H** = high/catastrophic, **M** = medium, **L** = low, — = not encountered.

#	Theme	P1	P2	P3	Root location
1	Recommender entity modelled per-position	H	H	—	database.py:135
2	No bulk operations across rows	H	H	M	pages/3_Recommenders.py
3	Vocabulary hardcoded in config.py	M	M	M	config.py:45-322
4	Schema-UI gap on positions table	H	L	L	database.py:77-94
5	Cascade asymmetry on delete_interview	—	—	H	database.py:707-719
6	No materials versioning / attachments	L	H	L	schema by design
7	Threshold tuning requires Python edit	L	M	L	config.py:320-322
8	Recommender email column missing	H	L	—	database.py:135-148
9	No responsive layout (tablet broken)	H	—	—	app.py:19; ui.py
10	No keyboard shortcuts	—	L	M	ui.py:1265 shield
11	Pipeline doesn't fit non-postdoc cycles	—	M	H	config.py:45-53,265
12	Reminder mailto per-position not per-recommender	L	M	—	pages/3_Recommenders.py:230
13	Dashboard date column has no year	M	L	—	app.py:378
14	No in-app backup affordance	L	H	—	pages/4_Export.py
15	init_db() on every page load (~500ms)	—	—	L	app.py:22

8. Friction Severity Heatmap

The same fifteen themes rendered as a heatmap make the persona-reach pattern visible at a glance. Three themes are universal (cells lit in all three columns); five are persona-specific. The pattern is informative for prioritization: **universal frictions are P0 candidates**; persona-specific high-severity frictions (P1 mobile, P2 versioning, P3 cascade defect) belong in P1/P2 tiers.



Figure 2. Friction severity matrix.

9. Verified Bug-class Findings

Four findings rise above UX friction and qualify as defects or feature-completion gaps. Each was verified against the source on the *feat/ui-redesign-v0.14.0* branch.

#	Finding	Severity	Location	Status
B1	R2 cascade asymmetry: add_interview promotes [APPLIED]→[INTERVIEW] database new has 18 reverse Confirmed. Position promoted via single (delete path)	High	pages/27-319	Confirmed
B2	Recommender email column missing from schema; Compose Reminder Email database by 135148 with an empty field. pages/3_Recommenders.py:142	High	pages/3_Recommenders.py:142	Confirmed
B3	Schema-UI gap on positions table: 10 nullable columns (location, single, mentor, phone, contact, portal_url, Copied, keywords, description, dead pages/1_Opportunities.py	High	pages/1_Opportunities.py	Confirmed
B4	Applications page calls upsert_application with applied_date; R1 cascade fires pages/2_Applications.py:506-510 vs database.py:506-510 Renowned	High	pages/2_Applications.py:506-510	Confirmed

Recommended response. File B1, B2, B3 as *fix*-typed commits before any new feature work; each is a few-hour change with high churn-risk payoff. B4 needs a one-line toast addition in *pages/2_Applications.py* matching the existing R2 toast pattern at *pages/2_Applications.py:645*.

10. Churn-risk Timeline + Retained Scope

The two charts below trace each persona's trust + active-usage index across the 12-week study, and the share of intended scope still active at the end of Month 3.




Figure 3. Trust + retained-usage index by week, with annotated near-quit events. The dashed line marks the qualitative quit threshold; the postdoc and humanities personas hover above it but do not cross.




Figure 4. Share of intended scope at end of Month 3 — split into still-used, covered-by-parallel-tool, and outright-abandoned slices.

Reading. P3 trends upward after Week 5 because the codebase rewards extension; his abandonment risk converts into contribution. P1 and P2 settle into a permanent ~50% trust plateau — the tool is retained but never owns the workflow. The retained-scope chart shows this directly: 20% of Aisha's intended workflow is outright abandoned (the Recommenders page and the tablet-on-bus tracking sessions).

11. What's Working — Strengths to Preserve

Six product strengths surfaced as positively-noted in two or more persona studies. These are not optional polish; they are the reasons each persona retained the tool through Month 3 despite the frictions above. Any feature work that risks degrading these properties should clear an explicit bar.

#	Strength	Personas	Why it matters
S1	Editorial-brutalist visual identity	P1, P2, P3	Only Streamlit app in surveyed competitor set with a distinct visual voice. P2: “the on
S2	Upcoming-deadlines panel + urgency D1, D2, D3	P1, P2, P3	The single feature that genuinely outperforms a spreadsheet. Daily-driver pane for all
S3	Markdown export on every write (D4)	P1, P2, P3	Builds durable trust that data survives the app itself. P3 commits the exported markdo
S4	Auto-promotion cascade (R1/R2/R3) P1 (love), P2 (yes), P3 (yes)	P1, P2, P3	Removes invisible work from the user's head when correct. When incorrect (see B1, B
S5	config.py import-time invariants	P3	Converted P3's near-fork into an upstream PR. The single feature that earned the code
S6	Quick-Add 6-field discipline (D6)	P1, P2, P3	Captures a new posting in genuinely under 30 seconds. Do not widen.

12. Competitive Landscape Analysis

Four competitive categories were surveyed against the Academic Application Tracker (AAT). For each category, two to three representative products were compared on pricing, deployment, and the features most relevant to AAT's academic-applicant use case. Full URLs are in the References appendix.

Category	Representative product	Pricing	Deployment	Notable
Academic-specific	Interfolio Dossier	\$48–60 / year	Cloud SaaS	Canonical letter-delivery service; mandatory account
	Academic Jobs Online	Free to applicant	Cloud SaaS	STEM-heavy; institution-side workflow
	Versatile PhD	Institutional subscription	Cloud SaaS	Content + community; not a tracker
General SaaS	Huntr	Free ≤40 jobs; \$10/mo Pro	Cloud + Chrome extension	Kanban + extension; no recommender model; inc
	Teal HQ	Free; \$29/mo Teal+	Cloud + extension	AI resume / cover-letter; ATS-keyword tuned
	JibberJobber	Free; \$9.95/mo Premium	Cloud SaaS	CRM-style; dated UI; generic
Notion / templates	Notion job-tracker template	Free template + account	Cloud SaaS	Fully customisable; no schema validation
	Etsy “Psych Postdoc Tracker” sheet	One-off (~\$5)	Google Sheets	Closest direct analog; explicit LOR + materials c
OSS self-host	JobSync (GitHub)	Free	Next.js + Docker	Generic; no academic vocab; AI resume review
	JobHunt (GitHub)	Free	React + ASP.NET 7	Aggregates scraped postings; generic

Positioning conclusion. The combination *local SQLite + academic vocabulary + first-class recommender workflow + markdown export + \$0 + no account* is not offered by any product surveyed. AAT does not replace Interfolio (still essential for delivering letters to committees). It replaces the **personal tracking spreadsheet that almost every academic applicant maintains alongside Interfolio**, with a real schema, deadline urgency view, interview-round modelling, and a portable markdown archive once the cycle ends.

13. Competitive Feature Matrix

Fourteen features × six products. ● = full coverage, ■ = partial, — = absent. AAT's column is densest in academic vocabulary, recommender workflow, materials matrix, status pipeline, urgency view, local data ownership, markdown export, and pricing — while leaving letter-delivery, mobile, browser extension, AI tailoring, and multi-device sync entirely uncovered.




Figure 5. Competitive feature matrix.

For marketing. The empty cells on the right half of AAT's column are not weaknesses to apologize for — they are the explicit non-goals stated in DESIGN §1.3 (no auth, no cloud, no mobile-first, no email/calendar, no multi-user). The right user is one for whom that list of absences is itself a selling point.

14. Prioritized Recommendations — P0 (churn-blockers)

Four P0 items should ship before any new feature work. Each addresses a verified bug or a high-severity friction encountered by at least two personas, and each is scoped to a few hours to a few days of engineering effort.

ID	Item	Effort	Impact	Closes
R0a	Promote recommenders to a global entity. New (id, name, email, relationship) + join table (position_id, rec_id, rec_email, asked_date, confirmed)	Multi-week	High	Frontier #6, Bug #2
R0b	Wire up the 10 unused positions columns into Quick-Add (location, source) and Overview (tag, mentor, bug_id, position_id, portal_url, stipend, key)	1-2 days	High	Frontier #4
R0c	Symmetric R2 cascade: delete_interview must retract status from [INTERVIEW] to [APPLIED] when Bug Deleted row was the only interview on	1 day	High	Bug #1
R0d	Add a recommender_email column and populate the mailto: To field. Trivial schema migration	1 day	Medium-High	Bug #2, Friction #8

15. P1 & P2 Recommendations + Effort-Impact Map

P1 — frequent friction (next minor).

ID	Item	Effort
R1a	Bulk operations: multi-select in Opportunities table; bulk status flip, requirement set, mark-submitted across Weeks.	Weeks.
R1b	In-UI Settings page for DEADLINE_ALERT_DAYS, RECOMMENDER_ALERT_DAYS, UPCOMING_WORK_DAYS_OPTIONS. Persist to a Half-day	Half-day
R1c	TRACKER_PROFILE switch (roadmap §12). Ship faculty + industry profiles alongside postdoc, each with its own STATUS_VALUES, INT Low-week	Low-week

P2 — quality of life (post-v1).

ID	Item	Effort
R2a	File attachments + per-document version metadata (roadmap §12).	Multi-week
R2b	Responsive layout — wrap Quick-Add columns and filter bars below ~900px.	Week
R2c	Keyboard shortcuts (/ search, n new, j/k row nav). The hotkey shield post-v0.14.0 already gates on single-char keys days	Days
R2d	In-app backup button: one-click copy postdoc.db to ~/Documents/aat-backup-YYYYMMDD.db.	Half-day
R2e	Grouped reminder mailto — single email lists all owed positions per recommender.	Half-day




Figure 6. Effort vs impact 2×2 for the twelve prioritized items.

16. Strategic Roadmap & Audience-specific Takeaways

Suggested four-tier shipping order.

v0.14.0 (current branch, before merge)

- B1, B2, B3 fixes (one fix commit each)

v0.15.0

- R0a Recommender entity refactor
- R0d email + mailto
- R0b schema-UI wireup

v0.16.0

- R1a Bulk operations
- R1b In-UI Settings page
- R2d Backup button
- R2e Grouped mailto

v0.17.0+

- R1c TRACKER_PROFILE switch (postdoc + faculty + industry)
- R2a File attachments + versioning
- R2b Responsive layout
- R2c Keyboard shortcuts

For the marketing manager.

Position AAT as **“the personal tracking layer beneath Interfolio”** — the missing piece of the academic applicant's stack, not a replacement for the letter-delivery service. Target users: PhD candidates and postdocs on the academic market, comfortable with *streamlit run*, who already maintain a tracking spreadsheet they're frustrated with. Lead with the editorial design (S1), the urgency-banded Upcoming panel (S2), and the markdown export safety net (S3). Lead away from competing on browser extensions, AI tailoring, or mobile-first UX — those are explicit non-goals (DESIGN §1.3) and chasing them would dilute the niche.

For the product team.

Two near-quit moments per persona, all in the first two weeks, all anchored to the same root cause: **data-entry tax on the recommender model**. R0a is the most leveraged single bet in the backlog. Ship it before any new feature. After that, R1c (TRACKER_PROFILE switch) is the move that converts the architecture's promise into delivered breadth — and the move that lets P3-style developer-evaluators contribute upstream instead of forking.

For the engineering team.

Four verified findings ready to land as *fix*: commits today. B1 and B4 are small (half-day each). B3 is a 1–2 day sweep with a pinning test that asserts every nullable column has a UI binding. B2 plus R0d together are approximately one day including a migration. The *config.py* invariant harness (*config.py*:123–148) is your friend for the larger R0a refactor; extend the invariants to cover the new join table on Day 1 of that work.

17. Appendix A — Verified Code References

Every load-bearing file:line reference cited in the body of this report was checked against the current source on branch *feat/ui-redesign-v0.14.0* (commit 2ded3ba). The table below records the verification result.

Reference	Claim	Verification result
database.py:135–148	recommenders schema (per-position, no email column)	Confirmed
database.py:77–94	10 nullable positions columns unused by UI	Confirmed
database.py:506–510	R1 cascade fires on applied_date NULL→non-NULL	Confirmed (P1 claim reclassified)
database.py:629	R2 cascade fires on add_interview	Confirmed
database.py:707–719	delete_interview has no reverse cascade	Confirmed (real defect)
pages/3_Recommenders.py:142	Mailto built with no To field	Confirmed
pages/3_Recommenders.py:240	build_compose_mailto callsite	Confirmed
pages/3_Recommenders.py:311	Recommender add path (per-position)	Confirmed
config.py:45–53	STATUS_VALUES pipeline (postdoc-shaped)	Confirmed
config.py:265	INTERVIEW_FORMATS = ['Phone','Video','Onsite','Other']	Confirmed
config.py:286–295	REQUIREMENT_DOCS hardcoded list	Confirmed
config.py:320–322	DEADLINE_ALERT_DAYS, RECOMMENDER_ALERT_DAYS	Confirmed
config.py:123–148	Import-time invariants on STATUS/COLORS/LABELS	Confirmed
app.py:336–410	Upcoming panel + urgency banding	Confirmed
app.py:19	layout='wide' set unconditionally	Confirmed
ui.py:1265	Hotkey shield (post-v0.14.0 narrowed to single-char)	Confirmed
DESIGN.md §1.3	Explicit non-goals (no auth, cloud, mobile-first, ...)	Confirmed
DESIGN.md §9.3	R1/R2/R3 cascade documentation	Confirmed

18. Appendix B — Methodology Limitations + References

Limitations.

- Simulated personas only. No live human users, no telemetry, no recorded sessions. Friction frequencies and trust-index trajectories are reasoned, not measured.
- Three personas are a deliberately small sample chosen for diversity (postdoc / faculty / industry-hybrid). A real study should add at least one mobile-first user, one international applicant, and one user with accessibility needs.
- Competitive analysis is desk research from public marketing pages and review sites; no hands-on trials of competitors were conducted.
- Severity ratings in the friction heatmap are composite expert judgements, not user-elicited Likert scores.
- The B1 defect (R2-delete asymmetry) was verified in source but not reproduced in a running app. A pinning test should accompany its fix.

References.

- [1] Interfolio Dossier — <https://www.interfolio.com/dossier/>
- [2] Interfolio pricing reference — <https://www.trustradius.com/products/interfolio-dossier/pricing>
- [3] Academic Jobs Online — <https://academicjobsonline.org/ajo>
- [4] Versatile PhD — <https://versatilephd.com/>
- [5] Huntr — <https://huntr.co>
- [6] Teal HQ — <https://www.tealhq.com>
- [7] Simplify — <https://simplify.jobs>
- [8] JibberJobber — <https://www.jibberjobber.com/pricing.php>
- [9] Notion job-applications template — <https://www.notion.com/templates/job-applications>
- [10] Etsy Psych Postdoc Tracker — <https://www.etsy.com/listing/4310751919/psych-postdoc-application-tracker>
- [11] Job-tracker comparison (Prentus) — <https://prentus.com/blog/we-found-the-5-best-job-tracker-tools-on-the-market>
- [12] JobSync (GitHub) — <https://github.com/Gsync/jobsync>
- [13] JobHunt (GitHub) — <https://github.com/jamerst/JobHunt>
- [14] AAT source — <https://github.com/YuZh98/academic-application-tracker> (branch feat/ui-redesign-v0.14.0 @ 2ded3ba)
- [15] Keep a Changelog — <https://keepachangelog.com/en/1.1.0/>
- [16] Semantic Versioning — <https://semver.org/spec/v2.0.0.html>

End of report. Regenerate with `.venv/bin/python scripts/build_ux_report.py`